Catalog

NVIDIA Training Course Catalog

August 2025

Introduction

NVIDIA offers training for diverse needs, giving individuals and teams across organizations what they need to advance their knowledge in AI, accelerated computing, data science, data center administration, graphics and simulation, networking, and more.

With access to high-performance computing, you'll learn how to train, optimize, and deploy neural networks using the latest deep learning tools, frameworks, and SDKs. You'll also learn how to assess, parallelize, optimize, and deploy NVIDIA-accelerated computing applications.

Our training program offers both self-paced online courses and instructor-led, prescheduled workshops. The self-paced courses range from ten minutes to eight hours and guide you through applying a specific technology, setting up a project, or administering solutions in a data center. Instructor-led workshops and boot camps go deeper into topic areas, teaching you how to implement a project or solution from end to end. Both types of courses give you valuable hands-on experience using the latest technologies.

Why Choose NVIDIA for Training?

- Learn how to build deep learning and accelerated computing applications for industries such as healthcare, robotics, autonomous driving, manufacturing, and more.
- Gain hands-on experience with the most widely used, industry-standard platforms including software, hardware, tools, and frameworks. Each student will have access to a fully configured, GPU-accelerated server in the cloud or access to NVIDIA solutions in our training lab.
- Become proficient in administering NVIDIA hardware and software solutions such as DGX™, InfiniBand, Cumulus, NVIDIA AI Enterprise, and more.
- Access instructor-led workshops and online courses from anywhere using just a laptop and internet connection.
- Acquire real-world expertise through content designed in collaboration with industry leaders such as Children's Hospital of Los Angeles, Mayo Clinic, and PwC.
- Earn NVIDIA certifications and course completion certificates to indicate subject matter competency and support your career growth.

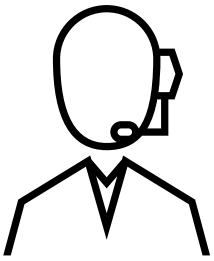

For team training, contact an NVIDIA training advisor, who will work with you to create a customized plan that addresses your team's specific training needs and is aligned to your business objectives and priorities.

Table of Contents

Instructor-Led Workshops for Developers

Accelerated Computing

Accelerating CUDA C++ Applications With Multiple GPUs 8
Fundamentals of Accelerated Computing with CUDA Python 8
Fundamentals of Accelerated Computing With Modern CUDA C++ 8
Fundamentals of Accelerated Computing with OpenACC 8
Scaling CUDA C++ Applications to Multiple Nodes 8

Data Science

Building AI-Based Cybersecurity Pipelines 8
Enhancing Data Science Outcomes With Efficient Workflow 8
Fundamentals of Accelerated Data Science 9

Deep Learning

Applications of AI for Anomaly Detection 9
Applications of AI for Predictive Maintenance 9
Bootstrapping Computer Vision Models With Synthetic Data 9
Building Conversational AI Applications 9
Building Transformer-Based Natural Language Processing Applications 9
Computer Vision for Industrial Inspection 9
Data Parallelism: How to Train Deep Learning Models on Multiple GPUs 10
Fundamentals of Deep Learning 10
Getting Started With AI on Jetson Nano 10
Model Parallelism: Building and Deploying Large Neural Networks 10

Generative AI and Large Language Models (LLMs)

Adding Knowledge to LLMs 10
Building Agentic AI Applications With LLMs 10
Building AI Agents With Multimodal Models 10
Building LLM Applications With Prompt Engineering 10
Building RAG Agents With LLMs 11
Deploying RAG Pipelines for Production at Scale 11
Generative AI With Diffusion Models 11
Rapid Application Development With Large Language Models (LLMs) 11

Online, Self-Paced Courses for Developers

Accelerated Computing

Accelerating CUDA C++ Applications With Concurrent Streams 12
An Even Easier Introduction to CUDA 12
Fundamentals of Accelerated Computing With CUDA Python 12
Fundamentals of Accelerated Computing With OpenACC 12
Getting Started With Accelerated Computing in Modern CUDA C++ 12
GPU Acceleration With the C++ Standard Library 12
Optimizing CUDA Machine Learning Codes With Nsight Profiling Tools 12
Scaling GPU-Accelerated Applications With the C++ Standard Library 13
Scaling Workloads Across Multiple GPUs With CUDA C++ 13

Data Science

Accelerate Data Science and Leverage Foundation Models in Digital Biology 13
Accelerate Data Science Workflows With Zero Code Changes 13
Accelerating Clustering Algorithms to Achieve the Highest Performance 13
Accelerating End-to-End Data Science Workflows 13
Analyzing and Visualizing Large Data Interactively Using Accelerated Computing 13
Digital Fingerprinting With Morpheus 14
Exploring Adversarial Machine Learning 14
RAPIDS Accelerator for Apache Spark 14
Structure From Chaos: Accelerate GraphRAG With cuGraph and NVIDIA NIM 14

Deep Learning

Applying AI Weather Models With NVIDIA Earth-2 14
Build High-Performance and AI-Enabled Sensor Processing Applications 14
Building a Brain in 10 Minutes 14
Building Real-Time Video AI Applications 14
Decentralized AI at Scale With NVIDIA FLARE 15
Disaster Risk Monitoring Using Satellite Imagery 15
Get Started With Highly Accurate Custom ASR for Speech AI 15
Getting Started With AI on Jetson Nano 15
Getting Started With Deep Learning 15
Introduction to Federated Learning With NVIDIA FLARE 15
Introduction to Graph Neural Networks 15
Medical AI Development With MONAI: Interactive Annotation Using NVIDIA NIM Microservices 15

Generative AI and Large Language Models (LLMs)

Augment Your LLM Using Retrieval-Augmented Generation 16
Building LLM Applications With Prompt Engineering 16
Building RAG Agents With LLMs 16

Evaluating RAG and Semantic Search Systems 16
Evaluation and Light Customization of Large Language Models 16
Generative AI Explained 16
Generative AI With Diffusion Models 16
Introduction to Deploying RAG Pipelines for Production at Scale 17
Introduction to Multimodal Data Curation 17
Introduction to NVIDIA NIM Microservices 17
Introduction to Transformer-Based Natural Language Processing 17
Rapid Application Development With Large Language Models (LLMs) 17
Sizing LLM Inference Systems 17
Streamlining Drug Discovery With NVIDIA BioNeMo Blueprints and NIM Microservices 17
Synthetic Tabular Data Generation Using Transformers 18
Techniques for Improving the Effectiveness of RAG Systems 18

Graphics and Simulation

Building a 3D Product Configurator With OpenUSD and Omniverse 18
Building and Deploying Digital Twin Applications With Omniverse Kit App Streaming 18
Creating and Customizing an Omniverse Extension 18
Develop, Customize, and Publish in Omniverse With Extensions 18
Developing an AI Background Generator With NIM 18
Developing an Omniverse Kit-Based Application 19
Fundamentals of Working With OpenUSD 19
How to Build a Native OpenUSD XR Application 19
How to Build OpenUSD Applications for Industrial Digital Twins 19
Learn OpenUSD: An Introduction to Strength Ordering 19
Learn OpenUSD: Asset Structure Principles and Content Aggregation 19
Learn OpenUSD: Creating Composition Arcs 19
Learn OpenUSD: Developing Data Exchange Pipelines 20
Learn OpenUSD: Learning About Stages, Prims, and Attributes 20
Learn OpenUSD: Setting Up Basic Animations 20
Learn OpenUSD: Traversing Stages 20
Learn OpenUSD: Understanding Model Kinds 20
Learn OpenUSD: Using Attributes 20
Learn OpenUSD: Working With Prims and Default Schemas 20
Synthetic Data Generation for Perception Model Training in Isaac Sim 20
Synthetic Data Generation for Training Computer Vision Models 21

Robotics

A Beginner's Guide to Autonomous Robots 21
Accelerating ROS 2 With NVIDIA GPU-Powered Libraries and AI Models 21
An Introduction to AI-Based Robot Development With Isaac ROS 21
An Introduction to NVIDIA Cosmos for Physical AI 21

An Introduction to Robot Learning and Isaac Lab 21
Develop, Simulate, and Deploy Robot Intelligence With Scaled Foundations 21
Developing Robots With Software-in-the-Loop (SIL) in Isaac Sim 21
Generating High-Quality Motion Data for Robotics With MobilityGen 22
Getting Started: Simulating Your First Robot in Isaac Sim 22
Ingesting Robot Assets and Simulating Your Robot in Isaac Sim 22
Leveraging ROS 2 and Hardware-in-the-Loop (HIL) in Isaac Sim 22
Software-in-the-Loop Testing for Robots With OpenUSD, Isaac Sim, and ROS 22
Transferring Robot Learning Policies From Simulation to Reality 22
Train Your First Robot in Isaac Lab 22
Train Your Second Robot in Isaac Lab 22

Instructor-Led Workshops for Administrators

AI and Data Science

AI Infrastructure and Operations: Professional Public Training 23
AI Infrastructure Professional: Public Training 23
AI Operations Professional: Public Training 23
NVIDIA AI Enterprise Deployment on Bare Metal Kubernetes Public Training 23

Cluster Administration

NVIDIA Base Command Manager 23

Ethernet Cumulus

Spectrum-X Networking Platform Administration - Private 24
Spectrum-X Networking Platform Administration - Public 24
NVIDIA Cumulus Linux: Public Bootcamp 24
NVIDIA Cumulus: Private Workshop 24
NVIDIA Cumulus Linux: Customized Advanced Training 24

InfiniBand

InfiniBand: Customized Course 25
InfiniBand: Professional Customized Training 25

NVIDIA DGX

DGX Onboarding for AI Practioners: Private Workshop 25
NVIDIA DGX H200/H100/A100 Administration: Private Workshop 25
NVIDIA DGX H200/H100/A100 Administration: Public Workshop 25
NVIDIA DGX BasePOD Administration: Private Workshop 26
NVIDIA DGX SuperPOD Administration: Private Workshop 26

Virtualization

NVIDIA AI Enterprise Administration: Public Bootcamp 26

Online, Self-Paced Courses for Administrators

AI and Data Science

AI for All—From Basics to Gen AI Practice 27

AI Infrastructure and Operations Fundamentals 27

NVIDIA AI Enterprise Administration for VMware vSphere Deployment 27

NVIDIA AI Enterprise for Azure Professionals 27

Cluster Administration

Base Command Manager Administration 27

Ethernet

Network Administration With the NVIDIA Onyx Switch System 28

RDMA Over Converged Ethernet (RoCE) From A to Z 28

InfiniBand

InfiniBand Essentials 28

InfiniBand Professional 28

Management

Data Center Management Made Easy With NVIDIA UFM 28

NVIDIA License System 29

Network

Ansible Essentials for Network Engineers 29

Introduction to Networking 29

MLXlink and MLXcables Debug Tools 29

NVIDIA BlueField DPU Administration 29

RDMA

The Fundamentals of RDMA Programming 29

Certifications

NVIDIA-Certified Associate: AI Infrastructure and Operations 30

NVIDIA-Certified Associate: Generative AI Large Language Models 30

NVIDIA-Certified Associate: Generative AI Multimodal 30

NVIDIA-Certified Professional: Accelerated Data Science 30

NVIDIA-Certified Professional: AI Infrastructure 30

NVIDIA-Certified Professional: AI Networking 30

NVIDIA-Certified Professional: AI Operations 31

NVIDIA-Certified Professional: OpenUSD Development 31

Instructor-Led Workshops for Developers

Workshop Name	Description			
Accelerated Computing				
Accelerating CUDA C++ Applications With Multiple GPUs	Use the NVIDIA Nsight Systems Visual Profile timeline to observe improvement opportunities and the impact of the techniques you learn. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	8 hours	$500	Yes
Fundamentals of Accelerated Computing with CUDA Python	Explore how to use Numba—the just-in-time, type-specializing Python function compiler—to create and launch CUDA kernels to accelerate Python programs on massively parallel NVIDIA GPUs. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese, Traditional Chinese	8 hours	$500	Yes
Fundamentals of Accelerated Computing With Modern CUDA C++	Learn how to write, compile, and run GPU-accelerated code, leverage CUDA core libraries to harness the power of massive parallelism provided by modern GPU accelerators, optimize memory migration between CPU and GPU, and implement your own algorithms. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Japanese, Korean, Simplified Chinese, Traditional Chinese	8 hours	$500	Yes
Fundamentals of Accelerated Computing with OpenACC	Learn the basics of OpenACC, a high-level language for programming on GPUs. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Japanese	8 hours	$500	Yes
Scaling CUDA C++ Applications to Multiple Nodes	Learn the tools and techniques needed to write CUDA C++ applications that can scale efficiently to clusters of NVIDIA GPUs. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$500	Yes
Data Science				
Building AI-Based Cybersecurity Pipelines	Build production-ready AI-based cybersecurity pipelines using NVIDIA Morpheus, a GPU-accelerated framework designed for real-time cybersecurity analytics. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$500	Yes
Enhancing Data Science Outcomes With Efficient Workflow	Learn how to develop and deploy an accelerated end-to-end data processing pipeline for large datasets. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$500	Yes

Back

Workshop Name	Description			
Fundamentals of Accelerated Data Science	Learn how to use GPU-accelerated tools to conduct data science faster, leading to more scalable, reliable, and cost-effective results. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$500	Yes

Deep Learning

Applications of AI for Anomaly Detection	Prepare data and build, train, and evaluate models using XGBoost, autoencoders, and GAN. Learn to detect anomalies in datasets with both labeled and unlabeled data as well as how to classify anomalies into multiple categories regardless of whether the original data was labeled. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$500	Yes
Applications of AI for Predictive Maintenance	Learn how to identify anomalies and failures in time series data, estimate the remaining useful life of the corresponding parts, and use this information to map anomalies to failure conditions. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Japanese, Traditional Chinese	8 hours	$500	Yes
Bootstrapping Computer Vision Models With Synthetic Data	In this course, you'll use NVIDIA Omniverse Replicator and the Omniverse Defect Extension to generate synthetic data. Next, you'll iterate on the dataset to train a deep neural network (DNN) to find target objects (scratches) in a scene. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$500	Yes
Building Conversational AI Applications	Learn how to quickly build and deploy a conversational AI pipeline including transcription, NLP, and speech. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	9 hours	$500	Yes
Building Transformer-Based Natural Language Processing Applications	Learn how to apply and fine-tune a Transformer-based deep learning model to natural language processing (NLP) tasks. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Japanese, Korean, Simplified Chinese, Traditional Chinese	8 hours	$500	Yes
Computer Vision for Industrial Inspection	Extract meaningful insights from the provided dataset using Pandas DataFrame. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese, Traditional Chinese	8 hours	$500	Yes

Back

Workshop Name	Description			
Data Parallelism: How to Train Deep Learning Models on Multiple GPUs	Learn how to distribute data across multiple GPUs during deep learning model training to make possible an incredible wealth of new applications utilizing deep learning. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$500	Yes
Fundamentals of Deep Learning	Learn the fundamental techniques and tools required to train a deep learning model. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Japanese, Korean, Simplified Chinese, Traditional Chinese	8 hours	$500	Yes
Getting Started With AI on Jetson Nano	Build and train a classification dataset and model with NVIDIA Jetson Nano. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Japanese, Korean, Simplified Chinese, Traditional Chinese	8 hours	$500	Yes
Model Parallelism: Building and Deploying Large Neural Networks	Learn how to train neural networks across multiple servers. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	8 hours	$500	Yes

Generative AI and Large Language Models (LLMs)

Adding Knowledge to LLMs	Transform open-source LLMs into domain-specialized AI assets. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$500	Yes
Building Agentic AI Applications With LLMs	Learn how to design intelligent agents that can be adapted for arbitrary environments at scale. Utilizing frameworks like LangGraph and NVIDIA NIM, you'll build agents capable of deep thought, long-horizon reasoning, content management, and real-time operations. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	8 hours	$500	Yes
Building AI Agents With Multimodal Models	Learn how to build neural network agents that reason across multiple data types using advanced fusion techniques, OCR, and NVIDIA AI Blueprints for real-world applications like robotics and healthcare. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$500	Yes
Building LLM Applications With Prompt Engineering	Learn how to build a variety of LLM-based applications through the use of modern prompt engineering techniques. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	8 hours	$500	Yes

Back

Workshop Name	Description			
Building RAG Agents With LLMs	Agents powered by large language models (LLMs) have shown great retrieval capability for using tools, looking at documents, and planning their approaches. This course will show you how to deploy an agent system in practice with the flexibility to scale up your system to meet the demands of users and customers. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Japanese	8 hours	$500	Yes
Deploying RAG Pipelines for Production at Scale	Participants will gain hands-on experience deploying, monitoring, and scaling RAG pipelines with the NIM Operator and learn best practices for infrastructure optimization, performance monitoring, and handling high traffic volumes. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$500	Yes
Generative AI With Diffusion Models	In this course, learners will take a deeper dive into denoising diffusion models, which are a popular choice for text-to-image pipelines. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	8 hours	$500	Yes
Rapid Application Development With Large Language Models (LLMs)	Get started quickly in developing LLM-based applications by exploring the open-source ecosystem, including pretrained LLMs. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese, Traditional Chinese	9 hours	$500	Yes

Back

Online, Self-Paced Courses for Developers

Course Name	Description			
Accelerated Computing				
Accelerating CUDA C++ Applications With Concurrent Streams	The concurrent overlap of GPU computation and the transfer of memory to and from the GPU can drastically improve the performance of CUDA applications. In this workshop, you'll learn to utilize CUDA streams to perform copy/compute overlap in CUDA C++. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	4 hours	$30	N/A
An Even Easier Introduction to CUDA	An interactive accompaniment to Mark Harris's popular blog post "An Even Easier Introduction to CUDA." > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	1 hour	Free	N/A
Fundamentals of Accelerated Computing With CUDA Python	Explore how to use Numba—the just-in-time, type-specializing Python function compiler—to create and launch CUDA kernels to accelerate Python programs on massively parallel NVIDIA GPUs. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$90	Yes
Fundamentals of Accelerated Computing With OpenACC	Learn the basics of OpenACC, a high-level language for programming on GPUs. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$90	Yes
Getting Started With Accelerated Computing in Modern CUDA C++	Learn how to write, compile, and run GPU-accelerated code, leverage CUDA core libraries to harness the power of massive parallelism provided by modern GPU accelerators, optimize memory migration between CPU and GPU, and implement your own algorithms. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	8 hours	$90	Yes
GPU Acceleration With the C++ Standard Library	Learn to write simple, portable, parallel-first applications using only standard C++ language features that can be compiled without modification to take advantage of NVIDIA GPU-accelerated environments. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	$30	N/A
Optimizing CUDA Machine Learning Codes With Nsight Profiling Tools	Learn how to use two powerful NVIDIA developer tools: Nsight Systems and Nsight Compute. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	$30	N/A

Back

Course Name	Description			
Scaling GPU-Accelerated Applications With the C++ Standard Library	Learn to write simple, portable, parallel-first applications using only standard C++ language features that can be compiled without modification to take advantage of NVIDIA GPU-accelerated environments. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	$30	N/A
Scaling Workloads Across Multiple GPUs With CUDA C++	Learn how to build robust and efficient CUDA C++ applications that can leverage all available GPUs on a single node. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	4 hours	$30	N/A

Data Science

Accelerate Data Science and Leverage Foundation Models in Digital Biology	Learn how foundation models can reveal novel biological insights, while leveraging NVIDIA's platform to accelerate the entire workflow for fast, scalable results. This course is free for a limited time. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	Free	N/A
Accelerate Data Science Workflows With Zero Code Changes	In this course, you'll learn to use RAPIDS to speed up your CPU-based data science workflows. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	1 hour	Free	N/A
Accelerating Clustering Algorithms to Achieve the Highest Performance	Learn how clustering algorithms like k-means, DBSCAN, and HDBSCAN are used to uncover patterns in data and power real-world applications, while leveraging NVIDIA GPUs to accelerate the entire workflow for fast, scalable results. This course is free for a limited time. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	Free	N/A
Accelerating End-to-End Data Science Workflows	Data science is about using scientific methods, processes, algorithms, and systems to analyze and extract insights from data. It empowers organizations to turn data into a valuable resource, leading to smarter decision-making, improved operations, and enhanced customer experiences. In this workshop, you'll learn how to use GPU-accelerated tools to conduct data science faster, leading to more scalable, reliable, and cost-effective results. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$90	Yes
Analyzing and Visualizing Large Data Interactively Using Accelerated Computing	Learn how to build responsive, interactive dashboards for large datasets using GPU-accelerated Python libraries. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	Free	N/A

Back

Course Name	Description			
Digital Fingerprinting With Morpheus	Perform digital fingerprinting with the NVIDIA Morpheus AI cybersecurity platform. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	1 hour	Free	N/A
Exploring Adversarial Machine Learning	In this course, which is designed to be accessible to both data scientists and security practitioners, you'll explore the security risks and vulnerabilities that adopting machine learning might expose you to. You'll also explore the latest techniques and tools being used by attackers and build some of your own attacks. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$90	Yes
RAPIDS Accelerator for Apache Spark	In this course, you'll walk through the details of how to leverage NVIDIA GPUs with Apache Spark to see workloads accelerate and reduce cost. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	$30	N/A
Structure From Chaos: Accelerate GraphRAG With cuGraph and NVIDIA NIM	Learn how to integrate large language models (LLMs) with NVIDIA NIM and cuGraph to create cutting-edge, graph-based AI solutions for handling complex, interconnected data. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	$30	N/A

Deep Learning

Applying AI Weather Models With NVIDIA Earth-2	Explore state-of-the-art AI weather prediction models and learn how to integrate them into custom workflows. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	3 hours	Free	N/A
Build High-Performance and AI-Enabled Sensor Processing Applications	Learn how to build, deploy, and scale domain-agnostic AI-enabled sensor processing applications. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	3 hours	$30	N/A
Building a Brain in 10 Minutes	The notebook explores the biological inspiration for early neural networks. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	10 minutes	Free	N/A
Building Real-Time Video AI Applications	Learn the skills you need to enable real-time transformation of raw video data from widely deployed camera sensors into deep learning-based insights. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	8 hours	$90	Yes

Back

Course Name	Description			
Decentralized AI at Scale With NVIDIA FLARE	Learn how to implement privacy-preserving distributed machine learning with NVIDIA FLARE. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	4 hours	Free	N/A
Disaster Risk Monitoring Using Satellite Imagery	Deploy a deep learning model to automate disaster management use cases. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	Free	Yes
Get Started With Highly Accurate Custom ASR for Speech AI	Learn to build, train, fine-tune, and deploy a GPU-accelerated automatic speech recognition service with NVIDIA Riva that includes customized features. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	3 hours	$30	N/A
Getting Started With AI on Jetson Nano	Build and train a classification dataset and model with the NVIDIA Jetson Nano. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	Free	Yes
Getting Started With Deep Learning	Learn how deep learning works through hands-on exercises in computer vision and natural language processing. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$90	Yes
Introduction to Federated Learning With NVIDIA FLARE	NVIDIA FLARE is a domain-agnostic, open-source, and extensible Python SDK for federated learning. It allows researchers and scientists to adapt existing machine learning/deep learning workflows to a federated paradigm. Learn what federated learning is, why industries such as healthcare, life sciences, and financial services adopt federated learning, and how it can potentially solve AI development challenges related to data privacy, security, and regulatory compliance. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	Free	N/A
Introduction to Graph Neural Networks	Learn the basic concepts, implementations, and applications of graph neural networks (GNNs) with hands-on interactive activities so you can get started using GNN as a graph analysis tool. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	2 hours	$30	N/A
Medical AI Development With MONAI: Interactive Annotation Using NVIDIA NIM Microservices	Learn to build end-to-end medical AI workflows using the latest tools. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	4 hours	Free	N/A

Back

Course Name	Description			
Generative AI and Large Language Models (LLMs)				
Augment Your LLM Using Retrieval-Augmented Generation	In this introductory course, we'll provide a high-level overview of retrieval-augmented generation and how it improves generative AI (gen AI). > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Japanese, Simplified Chinese	1 hour	Free	N/A
Building LLM Applications With Prompt Engineering	Learn how to build a variety of LLM-based applications through the use of modern prompt engineering techniques. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	8 hours	$90	Yes
Building RAG Agents With LLMs	Agents powered by large language models (LLMs) have shown great retrieval capability for using tools, looking at documents, and planning their approaches. This course will show you how to deploy an agent system in practice with the flexibility to scale up your system to meet the demands of users and customers. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Japanese, Simplified Chinese	8 hours	Free	Yes
Evaluating RAG and Semantic Search Systems	Understand the importance and unique challenges of deploying RAG and semantic search solutions in enterprise settings, including the integration of legacy and modern retrieval systems. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	3 hours	$30	N/A
Evaluation and Light Customization of Large Language Models	Learn how to assess and adapt large language models for specialized enterprise tasks through systematic evaluation and parameter-efficient fine-tuning. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	Free	Yes
Generative AI Explained	In this no-coding" course, learn generative AI concepts and applications, as well as the challenges and opportunities in this exciting field. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	2 hours	Free	N/A
Generative AI With Diffusion Models	Take a deeper dive into denoising diffusion models, which are a popular choice for text-to-image pipelines, with applications in creative content generation, data augmentation, simulation and planning, anomaly detection, drug discovery, personalized recommendations, and more. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$90	Yes

Back

Course Name	Description			
Introduction to Deploying RAG Pipelines for Production at Scale	The course focuses on teaching production-level deployment of LLM applications, especially enterprise-grade deployment of RAG pipelines. It covers various aspects for an end-to-end deployment using Helm and NVIDIA NIM microservices. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	4 hours	$90	N/A
Introduction to Multimodal Data Curation	NeMo Curator is a GPU-accelerated data-curation tool that improves generative AI model accuracy by processing text, image, and video data at scale for training and customization. It also provides pre-built pipelines for generating synthetic data to customize and evaluate generative AI systems. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	1 hour	Free	N/A
Introduction to NVIDIA NIM Microservices	Learn how NIM enables the building, deploying, and scaling of AI applications. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	Free	N/A
Introduction to Transformer-Based Natural Language Processing	Learn how Transformers are used as the building blocks of modern large language models (LLMs). You'll then use these models for various NLP tasks, including text classification, named-entity recognition (NER), author attribution, and question answering. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	6 hours	$30	N/A
Rapid Application Development With Large Language Models (LLMs)	Get started quickly in developing LLM-based applications by exploring the open-sourced ecosystem, including pretrained LLMs. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	8 hours	$90	Yes
Sizing LLM Inference Systems	This course teaches AI practitioners to optimize and deploy large language models using NVIDIA NIM microservices. It covers techniques like streaming, prefill, decoding, tensor parallelism, and in-flight batching. Students learn to benchmark models, select inference hyperparameters, and ensure efficient scaling for real-world applications. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	3 hours	$30	N/A
Streamlining Drug Discovery With NVIDIA BioNeMo Blueprints and NIM Microservices	Explore how BioNeMo Blueprints and NIM microservices can accelerate drug development, improve prediction accuracy, and enable more efficient analysis of biochemical data, showcasing real-world examples of their transformative potential in drug discovery. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	Free	N/A

Back

Course Name	Description			
Synthetic Tabular Data Generation Using Transformers	In this course, you'll learn the end-to-end development workflow for generating synthetic data using Transformers, including data preprocessing, model pretraining, fine-tuning, inference, and evaluation. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	4 hours	$30	N/A
Techniques for Improving the Effectiveness of RAG Systems	Learn techniques that can take your RAG system from an interesting proof-of-concept to a serious asset. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	4 hours	$30	N/A

Graphics and Simulation

Building a 3D Product Configurator With OpenUSD and Omniverse	Develop an Omniverse application for interactive, customizable 3D product configurators. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	Free	N/A
Building and Deploying Digital Twin Applications With Omniverse Kit App Streaming	Gain hands-on experience creating, deploying, and extending 3D digital twin applications using NVIDIA Omniverse. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	1.5 hours	$30	N/A
Creating and Customizing an Omniverse Extension	In this hands-on lab, you'll create an extension and customize it to make an interactive user interface, and learn how to extract Omniverse application commands to code the extension's functionality. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	Free	N/A
Develop, Customize, and Publish in Omniverse With Extensions	Learn how to customize the Omniverse experience with extensions using Python code. Extensions can be used for a wide variety of modifications, from spawning objects with a button press to applying custom physics on selected objects. Optimize a workflow by copying commonly repeated operations into an extension or add a new way to manipulate objects in the UI. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	8 hours	Free	Yes
Developing an AI Background Generator With NIM	Supercharge your Omniverse Kit-based application with NVIDIA NIM microservices. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	Free	N/A

Back

Course Name	Description			
Developing an Omniverse Kit-Based Application	To build applications from scratch, NVIDIA offers Omniverse Kit SDK and free templates to build starter applications that can be easily customized and extended. This course provides the fundamentals of using Kit SDK to develop such applications. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	Free	N/A
Fundamentals of Working With OpenUSD	Learn how you can use OpenUSD for nondestructive workflows, how layers can help with ease and speed of scene composition, and how to use OpenUSD for data separation and reuse to accelerate 3D workflows in industrial use cases. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	Free	N/A
How to Build a Native OpenUSD XR Application	In this hands-on lab, learn how to take advantage of Universal Scene Description (OpenUSD) to accelerate your extended reality (XR) development and enhance visual fidelity like never before. This session will equip you with the skills and tools necessary to build, customize, and stream your own OpenUSD-native XR applications using NVIDIA Omniverse and NVIDIA CloudXR. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	Free	N/A
How to Build OpenUSD Applications for Industrial Digital Twins	The NVIDIA Omniverse platform provides developers with SDKs, APIs, and microservices for developing OpenUSD-based workflows and applications that enable industrial-scale 3D digital twins, from planning to simulation and operations. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	Free	N/A
Learn OpenUSD: An Introduction to Strength Ordering	Get introduced to fundamental terms in OpenUSD and get hands-on practice with their implementation in Python. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	4 hours	Free	N/A
Learn OpenUSD: Asset Structure Principles and Content Aggregation	Explore the fundamental principles of asset structuring and learn how to leverage best practices for your OpenUSD scene organization. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	3 hours	Free	N/A
Learn OpenUSD: Creating Composition Arcs	Explore the fundamental concepts of OpenUSD composition, including layers, sublayers, references, payloads, variant sets, inherits, and specializes. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	3 hours	Free	N/A

Back

Course Name	Description			
Learn OpenUSD: Developing Data Exchange Pipelines	Explore the fundamentals of OpenUSD data exchange, learn techniques for data extraction and transformation, and gain hands-on experience in asset validation. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	1.5 hours	Free	N/A
Learn OpenUSD: Learning About Stages, Prims, and Attributes	Get introduced to fundamental terms in OpenUSD and get hands-on practice with their implementation in Python. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	4 hours	Free	N/A
Learn OpenUSD: Setting Up Basic Animations	Get introduced to fundamental terms in OpenUSD and get hands-on practice with their implementation in Python. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	4 hours	Free	N/A
Learn OpenUSD: Traversing Stages	Get introduced to fundamental terms in OpenUSD and get hands-on practice with their implementation in Python. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	30 minutes	Free	N/A
Learn OpenUSD: Understanding Model Kinds	Get introduced to model kinds in OpenUSD and experiment with their implementation in Python. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	15 minutes	Free	N/A
Learn OpenUSD: Using Attributes	Get introduced to fundamental terms in OpenUSD and get hands-on practice with their implementation in Python. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	30 minutes	Free	N/A
Learn OpenUSD: Working With Prims and Default Schemas	Learn about what it means to create a prim without a schema, and review the pre-built, default schemas that ship with OpenUSD. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	45 minutes	Free	N/A
Synthetic Data Generation for Perception Model Training in Isaac Sim	Learn to train and deploy perception models using synthetic data generation (SDG), applying domain randomization and simulation for real-world robotics. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	2 hours	Free	N/A

Back

Course Name	Description			
Synthetic Data Generation for Training Computer Vision Models	This course uses a food manufacturing example to demonstrate the power of Omniverse Replicator synthetic data generation to train computer vision models. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	3 hours	$30	N/A

Robotics

A Beginner's Guide to Autonomous Robots	Get an introduction to autonomous robots and robotic systems. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	1 hour	Free	N/A
Accelerating ROS 2 With NVIDIA GPU-Powered Libraries and AI Models	Learn how to accelerate ROS 2 workloads using NVIDIA's latest GPU-powered libraries for AI and robotics. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	3 hours	$30	N/A
An Introduction to AI-Based Robot Development With Isaac ROS	Get an overview of NVIDIA Isaac, Isaac ROS, and NITROS. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	30 minutes	Free	N/A
An Introduction to NVIDIA Cosmos for Physical AI	Gain hands-on experience with Cosmos world foundation models (WFMs) and tools to generate data to train physical AI. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	$30	N/A
An Introduction to Robot Learning and Isaac Lab	Explore the fundamentals of robot learning with Isaac Lab. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	3 hours	Free	N/A
Develop, Simulate, and Deploy Robot Intelligence With Scaled Foundations	Learn end-to-end robot development using Scaled Foundations' browser-based Open GRID platform, powered by NVIDIA Isaac Sim. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	1 hour	Free	N/A
Developing Robots With Software-in-the-Loop (SIL) in Isaac Sim	Learn the fundamentals of software-in-the-loop (SIL) using NVIDIA Isaac Sim and ROS 2 to test, validate, and develop robotics software in simulated environments. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	2 hours	Free	N/A

Back

Course Name	Description			
Generating High-Quality Motion Data for Robotics With MobilityGen	Create diverse datasets for wheeled robots, quadrupeds, and humanoids using MobilityGen. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	1.5 hours	Free	N/A
Getting Started: Simulating Your First Robot in Isaac Sim	Build foundational skills in robotics simulation and control with Isaac Sim, the first step in the Isaac Sim Learning Path. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	1.5 hours	Free	N/A
Ingesting Robot Assets and Simulating Your Robot in Isaac Sim	Learn to import robotic assets, add sensors, and run simple simulations. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	1 hour	Free	N/A
Leveraging ROS 2 and Hardware-in-the-Loop (HIL) in Isaac Sim	Bridge the gap between simulation and real-world robotics using ROS 2, Isaac Sim, and NVIDIA Jetson with HIL. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English, Simplified Chinese	2 hours	Free	N/A
Software-in-the-Loop Testing for Robots With OpenUSD, Isaac Sim, and ROS	Assemble a virtual environment with robots, and simulate and validate their basic movements with Robot Operating System (ROS) commands. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	Free	N/A
Transferring Robot Learning Policies From Simulation to Reality	Learn the principles of developing effective robot learning policies. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	1 hour	Free	N/A
Train Your First Robot in Isaac Lab	Introducing reinforcement learning for robotics with Isaac Lab. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	3 hours	Free	N/A
Train Your Second Robot in Isaac Lab	Go further with reinforcement learning for robotics with Isaac Lab. > Learn More			
	Languages	**Duration**	**Price**	**Certificate**
	English	2 hours	Free	N/A

Back

Instructor-Led Workshops for Administrators

Workshop Name	Description			
AI and Data Science				
AI Infrastructure and Operations: Professional Public Training	Hands-on training course that explores configuration, management, and troubleshooting of AI infrastructure and operations. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	44 hours	$5,500 (excludes tax, if applicable)	Available
AI Infrastructure Professional: Public Training	Hands-on training course that explores optimizing efficiency, reliability, and scalability for deploying AI environments. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	28 hours	$3,000 (excludes tax, if applicable)	Available
AI Operations Professional: Public Training	Hands-on training course that explores operating AI data centers, including provisioning and management, running AI workloads, and implementing AI virtualization. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	24 hours	$3,000 (excludes tax, if applicable)	Available
NVIDIA AI Enterprise Deployment on Bare Metal Kubernetes Public Training	This hands-on training course explores architecture, installation, configuration, operation, and management of NVIDIA AI Enterprise. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	24 hours	$3,000 (excludes tax, if applicable)	N/A
Cluster Administration				
NVIDIA Base Command Manager	This course provides an overview of NVIDIA Base Command Manager, including managing nodes and software images, monitoring devices and jobs, managing users, and configuring workload managers. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	12 hours	Contact us	N/A

Back

Workshop Name	Description			
Ethernet Cumulus				
Spectrum-X Networking Platform Administration - Private	In this course, participants will gain essential skills to deploy and manage AI clusters using the NVIDIA Spectrum-X networking platform. Through instructor-led sessions and hands-on labs in the NVIDIA Air environment, participants will learn to provision and manage AI clusters. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	12 hours	Contact us	N/A
Spectrum-X Networking Platform Administration - Public	In this course, participants will gain essential skills to deploy and manage AI clusters using the NVIDIA Spectrum-X networking platform. Through instructor-led sessions and hands-on labs in the NVIDIA Air environment, participants will learn to provision and manage AI clusters. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	12 hours	$1,500 (excludes tax, if applicable)	N/A
NVIDIA Cumulus Linux: Public Bootcamp	Learn how to install, deploy, configure, and troubleshoot Cumulus-based networks. This course offers a perfect blend of hands-on training and theoretical education. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	12 hours	$1,500 (excludes tax, if applicable)	Available
NVIDIA Cumulus: Private Workshop	In this hands-on private training, you'll learn about NVIDIA Cumulus OS architecture and installation, configuration, operation, and management of Cumulus Linux running on NVIDIA switches. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	20 hours	Contact us	Available
NVIDIA Cumulus Linux: Customized Advanced Training	This course focuses on how to build and operate a state-of-the-art data center or storage fabric with emphasis on troubleshooting. The course covers advanced topics such as filtering, quality of service (QoS), Ethernet VPN multihoming (EVPN-MH), monitoring, and active testing. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	12 hours	Contact us	N/A

Back

Workshop Name	Description			
InfiniBand				
InfiniBand: Customized Course	In this course, you'll learn about InfiniBand architecture and how to manage, monitor, and troubleshoot your InfiniBand network. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	16 hours	Contact us	Available
InfiniBand: Professional Customized Training	In this course, you'll learn about InfiniBand and NVIDIA Cumulus architecture and how to manage, monitor, and troubleshoot triad deployment-based networks. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	16 hours	Contact us	Available
NVIDIA DGX				
DGX Onboarding for AI Practioners: Private Workshop	This course introduces users to NVIDIA DGX systems and NVIDIA AI Enterprise software. Participants will learn how these two technologies can be combined to accelerate the development and deployment of AI applications. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	8 hours	Contact us	N/A
NVIDIA DGX H200/H100/ A100 Administration: Private Workshop	This course provides an overview of NVIDIA DGX systems, tools for in-band and out-of-band management, NVIDIA NGC, the basics of running workloads, and specific management tools and command-line interface (CLI) commands. In addition, this course includes content on Multi-Instance GPU (MIG), managing storage, performance validation, and other system management tools and concepts. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	16 hours	Contact us	N/A
NVIDIA DGX H200/H100/ A100 Administration: Public Workshop	This course provides an overview of DGX systems and tools for in-band and out-of-band management, the basics of running workloads, specific management tools, and CLI commands. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	16 hours	$1,500 (excludes tax, if applicable)	N/A

Back

Workshop Name	Description			
NVIDIA DGX BasePOD Administration: Private Workshop	This course provides an overview of DGX BasePOD components and related processes, including the NVIDIA DGX A100 system, InfiniBand and Ethernet networks, tools for in-band and out-of-band management, NGC, the basics of running workloads, and specific management tools and CLI commands. It includes instructions for managing vendor-specific storage per the architecture of your specific DGX BasePOD solution. **>** Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	16 hours	Contact us	N/A
NVIDIA DGX SuperPOD Administration: Private Workshop	This course is designed to help IT professionals successfully administer all aspects of a DGX SuperPOD cluster, including compute, storage, and networking. **>** Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	16 hours	Contact us	N/A

Virtualization

NVIDIA AI Enterprise Administration: Public Bootcamp	This course covers the platform and solution overview, hardware and software architecture, deployment options, licensing, temporal and spatial GPU partitioning, scaling, comprehensive validation, management, maintenance, monitoring, and troubleshooting. **>** Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	12 hours	$1,500 (excludes tax, if applicable)	N/A

Back

Online, Self-Paced Courses for Administrators

Workshop Name	Description			
AI and Data Science				
AI for All—From Basics to Gen AI Practice	This introductory course provides invaluable insights into the evolving landscape of AI. Whether you're a seasoned professional or just beginning your journey into AI, this course is essential for staying ahead in today's rapidly evolving technological landscape. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	2.5 hours	Free	N/A
AI Infrastructure and Operations Fundamentals	In this course, we'll start with an introduction to AI, where we'll cover basic AI concepts and principles. Then we'll delve into data center and cloud infrastructure before exploring AI operations. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	7 hours	$50 (excludes tax, if applicable)	Available
NVIDIA AI Enterprise Administration for VMware vSphere Deployment	This course covers the platform and solution overview, hardware and software architecture, deployment options, licensing, temporal and spatial GPU partitioning, scaling, comprehensive validation, management, maintenance, monitoring, and troubleshooting. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	8 hours	$100 (excludes tax, if applicable)	N/A
NVIDIA AI Enterprise for Azure Professionals	The NVIDIA AI Enterprise for Azure IT Professionals training is a comprehensive, hands-on course designed to equip Azure professionals with the skills needed to deploy, manage, and validate AI workloads in production environments using the NVIDIA AI Enterprise solution. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	9 hours	$100 (excludes tax, if applicable)	N/A
Cluster Administration				
Base Command Manager Administration	This course is based on NVIDIA Base Command Manager and gives an overview of the cluster management tools, Bright View, and cluster management shell (CMSH). > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	5 hours	Free	N/A

Back

Workshop Name	Description			
Ethernet				
Network Administration With the NVIDIA Onyx Switch System	This course provides the required set of skills to configure and manage NVIDIA Ethernet switch systems. You'll learn in depth layer 2 configurations such as virtual local area network (VLAN), Spanning Tree Protocol (STP), link aggregation (LAG), and multi-chassis link aggregation (MLAG), as well as how to configure layer 3 features such as Border Gateway Protocol (BGP). > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	3 hours	$100 (excludes tax, if applicable)	N/A
RDMA Over Converged Ethernet (RoCE) From A to Z	In this course, you'll learn what RoCE is, how it works, the different network types RoCE can run over, and how to configure RoCE for each network type. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	2 hours	Free	N/A
InfiniBand				
InfiniBand Essentials	This self-paced course covers the fundamental first steps into the world of InfiniBand. If you're looking to become more familiar with InfiniBand's benefits, uses, architecture layers, and management concepts, this is the best place to start. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	1.5 hours	Free	N/A
InfiniBand Professional	This course covers the fundamentals of the InfiniBand technology from a usability point of view and builds on the details of the InfiniBand architecture specification. You'll learn how to install, configure, manage, troubleshoot, and monitor your InfiniBand network. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	6 hours	$200 (excludes tax, if applicable)	Available
Management				
Data Center Management Made Easy With NVIDIA UFM	Learn about NVIDIA Unified Fabric Manager (UFM) and its capabilities, advantages, and components through a set of interactive learning units, videos, and simulators. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	3 hours	$50 (excludes tax, if applicable)	N/A

Back

Workshop Name	Description			
NVIDIA License System	NVIDIA License System (NLS) is a new licensing solution that supports the continued expansion of the NVIDIA enterprise software portfolio. This course will help you learn about NLS and how you can move from your existing licensing solution to NLS. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	2 hours	Free	N/A

Network

Ansible Essentials for Network Engineers	In this course, you'll explore a variety of Ansible modules and write playbooks specifically adapted to modern data centers. This course includes an exclusive hands-on lab environment and exercises to practice real-world scenarios in real cloud environments. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	3 hours	$50 (excludes tax, if applicable)	N/A
Introduction to Networking	In this course, we'll cover the basics of Ethernet technology and understand how data is forwarded in an Ethernet network. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	1 hour	Free	N/A
MLXlink and MLXcables Debug Tools	In this course, you'll learn about the MLXlink and MLXcables debug tools. These debug tools are used for both basic link troubleshooting and for analyzing more complex link characteristics. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	2 hours	Free	N/A
NVIDIA BlueField DPU Administration	Learn the basic concepts of BlueField DPUs as a platform for accelerated data center computing. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	3 hours	$50 (excludes tax, if applicable)	N/A

RDMA

The Fundamentals of RDMA Programming	This course allows C programmers to dive into the remote direct-memory access (RDMA) programming world without previous experience in networking or RDMA programming. We've also added tips and tricks, as well as do's and don'ts, so the skills you acquire will truly serve you when you need them. > Learn More			
	Languages	**Duration**	**Price**	**Certification Exam**
	English	4 hours	$50 (excludes tax, if applicable)	N/A

Back

Certifications

Certification Name	Description		
NVIDIA-Certified Associate: AI Infrastructure and Operations	This is an entry-level credential that validates the foundational concepts of AI computing related to infrastructure and operations. The exam is online and proctored remotely, includes 50 questions, and has a 60-minute time limit. > Learn More		
	Languages	**Duration**	**Price**
	English	1 hour	$135 (excludes tax, if applicable)
NVIDIA-Certified Associate: Generative AI Large Language Models	An entry-level credential that validates the foundational concepts for developing, integrating, and maintaining AI-driven applications using generative AI and large language models (LLMs) with NVIDIA solutions. The exam is online and proctored remotely, includes 50 questions, and has a 60-minute time limit. > Learn More		
	Languages	**Duration**	**Price**
	English	1 hour	$135 (excludes tax, if applicable)
NVIDIA-Certified Associate: Generative AI Multimodal	An entry-level credential that validates the foundational skills needed to design, implement, and manage AI systems that synthesize and interpret data across text, image, and audio modalities. The exam is online and proctored remotely, includes 50 questions, and has a 60-minute time limit. > Learn More		
	Languages	**Duration**	**Price**
	English	1 hour	$135 (excludes tax, if applicable)
NVIDIA-Certified Professional: Accelerated Data Science	This is an intermediate-level credential that validates a candidate's proficiency in leveraging GPU-accelerated tools and libraries for data science workflows. The exam is online and proctored remotely, includes 60–70 questions, and has a 90-minute time limit. > Learn More		
	Languages	**Duration**	**Price**
	English, Simplified Chinese	1.5 hours	$200 (excludes tax, if applicable)
NVIDIA-Certified Professional: AI Infrastructure	This is an intermediate-level credential that validates a candidate's ability to deploy, manage, and maintain AI infrastructure by NVIDIA. The exam is online and proctored remotely, includes 65 questions, and has a 90-minute time limit. > Learn More		
	Languages	**Duration**	**Price**
	English	1.5 hours	$400 (excludes tax, if applicable)
NVIDIA-Certified Professional: AI Networking	This is an intermediate-level credential that validates a candidate's ability to deploy and configure environments that leverage NVIDIA's advanced networking technologies. The exam is online and proctored remotely, includes ~65 questions, and has a 90-minute time limit. > Learn More		
	Languages	**Duration**	**Price**
	English	1.5 hours	$400 (excludes tax, if applicable)

Back

Certification Name	Description		
NVIDIA-Certified Professional: AI Operations	This is is an intermediate-level credential that validates a candidate's ability to monitor, troubleshoot, and optimize AI infrastructure by NVIDIA. The exam is online and proctored remotely, includes 50 questions, and has a 90-minute time limit. > Learn More		
	Languages	**Duration**	**Price**
	English	1.5 hours	$400 (excludes tax, if applicable)
NVIDIA-Certified Professional: OpenUSD Development	This is an intermediate-level credential that validates a candidate's ability to build, maintain, and optimize 3D content creation pipelines using OpenUSD. The exam is online and proctored remotely, includes 60–70 questions, and has a 90-minute time limit. > Learn More		
	Languages	**Duration**	**Price**
	English	1.5 hours	$200 (excludes tax, if applicable)

Ready to Get Started?

To get started with hands-on training, visit nvidia.com/en-us/learn/organizations

For questions, contact us.

© 2025 NVIDIA Corporation and affiliates. All rights reserved. NVIDIA, the NVIDIA logo, Base Command, BioNeMo, BlueField, CloudXR, Cosmos, CUDA, Cumulus, DGX, DGX BasePOD, DGX SuperPOD, FLARE, Isaac, Isaac Sim, Jetson, Jetson Nano, Nemo, NGC, NIM, Nsight, Omniverse, Onyx, RAPIDS, Riva, Triton Inference Server, and UFM are trademarks and/or registered trademarks of NVIDIA Corporation and its affiliates in the U.S. and other countries. Other company and product names may be trademarks of the respective companies with which they are associated. All other trademarks are property of their respective owners. 4050801. AUG25